

The Sybil Attack—Theory and Practice

Kelong Cong
Delft University of Technology
k.cong@student.tudelft.nl

ABSTRACT

The sybil attack is an attack where an entity can fully control numerous identities. It is one of the most important attacks because it leads to numerous consequences such as slandering and identity theft. In this work, we survey the practical and the theoretical aspects of the sybil attack. On the practical side, we demonstrate the severity of the sybil attack using real-world examples, it includes an experimental study to gain first-hand knowledge about the properties of the sybils. We then consider the sybil attack as an umbrella term for all the attacks that require the use of sybils and describe those attacks individually. The theoretical side of this work explains the underlying principals of the defence mechanisms and categorise them according to the main idea. From our findings, we realise that the sybil attack remains unsolved. Many of the defence mechanisms do not generalise well into scenarios in the real-world, and they break down when their assumptions are not satisfied. We hope this work serves as a cornerstone for the future sybil defence mechanisms.

1. INTRODUCTION

A recent article in the Atlantic describes how fake Twitter accounts are shaping the 2016 US presidential election [29]. Over a third of pro-Trump tweets and almost a fifth of pro-Clinton tweets, totalling at about 1 million, came from sybils¹. Even earlier, in the 2010 Massachusetts senate race, researchers found evidence that Republican campaigners also used fake accounts to manipulate Google real-time search results to tip the balance in their favour, effectively causing a spread of misinformation [55]. These are examples of the sybil attacks and they are a threat to democracy because opinions of real users are eclipsed by an overwhelming number of fake identities.

The sybil attack, first described by Douceur [20], is an attack where an entity can assume multiple identities, or sybils, and then attack either another entity or undermine the whole system. For example, a malicious Twitter user can create many fake identities and have them follow his real identity, thus creating a false sense reputation. It is one of the most important attacks because it leads to numerous consequences including but not limited to spreading false information, identity theft [4] and ballot stuffing [3]. Furthermore, to the best of our knowledge, there is no general solution for preventing the sybil attack.

There has been over a decade of work on the sybil attack

¹In this work we use “sybil” as a noun to mean a fake identity.




Figure 1: Screenshot of a commercial service for buying fake Facebook likes called boostlikes.com.

from both perspective—attackers and defenders. Many previous surveys focus only on the defender’s perspective or a particular class of defence mechanisms [49, 54, 61, 40]. The purpose of this work is to survey and provide a broad view on all aspects of the sybil attack. This work should be seen as an introductory material for readers who wish to familiarise themselves with research in this area.

This survey is organised as follows. First, we illustrate the importance of the sybil attack by looking at how researchers and black-hat hackers mounted the attack in the real-world; we also experiment with fake Twitter followers and look at their properties (section 2). Next, we give the models and definitions that we use in this work and some important theoretical results (section 3). With the fundamentals in place, we describe the various types of attacks (section 4) and defence mechanisms (section 5). Finally, we present the related work in section 6 and conclude in section 7.

2. REAL-WORLD ATTACKS

We begin our survey by showing some alarming sybil attacks happening in the real-world. First, we summarise a few amusing results from external studies, specifically on the convenience and the scale of the sybil attack. Next, we describe our experimental study on fake Twitter accounts.



Figure 2: Screenshot of the main banner on socialformulae.com.

2.1 External Studies

In the introduction we showed how Twitter is used to manipulate public opinion in elections. But this capability is not only accessible to campaigners with a large budget. There are marketplaces where anybody can purchase false reputation scores such as Twitter followers. Boost-Likes shown in Figure 1 is a professionally presented website, it offers a large range of services including Facebook likes, Twitter followers, Instagram followers and YouTube views. SocialFormulae (Figure 2) is a similar service but at a much lower price point, one thousand Twitter followers is only \$9.99.

SadBotTrue and its related website Socialpuncher publishes studies on social media fraud. Two of their studies are particularly useful for demonstrating the scale of the sybil attack and the obliviousness of Twitter. Firstly, there exist a sybil group that consist of 3 million accounts. Since their creation, they generated 2.6 billion tweets. Surprisingly, all the 3 million accounts were created on the same day—22/10/2013, and the account names are simply numbered sequentially [65]. Such an obvious activity should be easily detectable by Twitter, but these accounts are still not closed at the time of writing. Secondly, the top-100 Twitter users have 523 million unique followers between them, but 310 million are sybils, that is almost 60% [73]. Suppose the sybils all belong to the same attacker, then they can effectively suppress the opinions of the real users.

Clearly, the defence mechanisms employed by popular social network websites are not adequate to combat the sybil attack. If the sybils infiltrate even more of our cyberspace, then it may become a form of censorship. Effectively taking away our right to freedom of speech.

Speaking of censorship, around a million [60] people use Tor (The Onion Router) [19] to access the uncensored internet when living in authoritarian regimes such as China, or wish to uphold their privacy from illegal mass surveillance by intelligence agencies. Unfortunately, Tor suffered from a sybil attack. In January 2014, 115 bogus relays joined the Tor network. Six months later, it was discovered that those relays were using a protocol vulnerability to deanonymise users and find the location of hidden services. It is unclear to the Tor developers which users are affected or what information was retrieved, thus it is assumed that users who used Tor between that period are all affected [18]. In fact, Tor depends on the fact that majority of the relays are good

to guarantee anonymity with a high probability. If the network contains a large proportion of sybils, then users can be easily deanonymised.

2.2 Small Experiment

We show it takes less than one hour to obtain over 1000 sybils and investigate the properties of those sybils. Twitter is selected for two reasons, (1) it is one of the most popular social networks, (2) it is vulnerable to the sybil attack as described in subsection 2.1.

2.2.1 Buying Experience

Before we can purchase sybils, we needed our own accounts. One was created on the 25th of November 2016. We made sure our account has 0 tweets and is not advertised in any other medium over the duration of this experiment. This guarantees that all of its followers are sybils.

We purchased the “1,000 Followers” product from CoinCrack² for \$9. No information was asked except for an email address for them to deliver the receipt. Followers started coming in almost immediately after we made our purchase. This is especially interesting because we made the purchase using Bitcoin, but CoinCrack did not wait for 10 minutes for our transaction to be confirmed. No more than 1 hour later, our fresh account with 0 tweets has accumulated 1,300 followers. These followers are sybils, because CoinCrack’s website states “We broker followers in real time from dozens of engineers across the globe who specialize in creating/maintaining large amounts of Twitter profiles”.

2.2.2 Sybil Properties

Sybils themselves often form relationships in order to emulate real users. To find these relationships, we crawled our followers recursively for 72 hours using `tweetf0rm`³. We obtained 3 million nodes and 6 million edges at the end of the crawl. Every node represents an account and every edge represents a “following” relationship. For example, if *A* follows *B* then *A* has a directed edge to *B*. It is unlikely that all 3 million nodes are sybils. A reasonable assumption is that real users are unlikely to follow fake accounts because they do not create interesting or original content. Thus in this work we assume the majority of the 3 million accounts are sybils.

²<https://coincrack.com/>

³<https://github.com/bianjiang/tweetf0rm>




Figure 3: Graph of out-degree distribution for a graph with 3,248,093 nodes and 6,062,427 edges.

Faloutsos et al. [22] analysed the internet topology in 1999 using log-log plots where the x -axis is the out-degree and the y -axis is the frequency. We use the a similar methodology to analyse the Twitter graph. The result is shown in Figure 3. The regression line (green) is in the form of $p(x) = ax^k$, which is a power-law function. To determine the goodness-of-fit, we compute the coefficient of determination R^2 . R^2 is a value between 0 and 1, where 0 indicates no correlation and 1 indicates that the regression line perfectly fits the data. In this case $R^2 = 0.86$, which is an indication that the Twitter graph does have some properties of a power-law network. This is an interesting observation because many other networks such as the World-Wide Web, biological networks and other social networks also have the same property [22, 27]. We do not have a verifiable explanation for this property, but we speculate that it is because the attackers want to avoid detection, so they structure their sybil group in a way that look like a group of real users.

3. MODEL AND DEFINITION

In this section, we introduce the models and definitions used to describe this sybil attack, we use these in the remainder of the survey. Further, we give a few impossibility results to illustrate the difficulty of the sybil attack

3.1 Definitions

The sybil attack is coined by Douceur [20] in 2002 in the context of P2P (peer-to-peer) systems. But people were well aware of it before 2002. For instance in 2001, Friedman and Resnick used the term “cheap pseudonyms” to describe sybils [64].

Douceur defined the sybil attack as forging multiple *identities* under the same *entity* [20]. An entity can be for example a physical user of the system and identities are how entities present themselves to the system. Thus, a local entity has no direct knowledge of remote entities, only their identities. The forged identities do not necessarily follow the protocol specified by the underlying network, i.e. they assume the characteristics of Byzantine fault [43]. In this work, we use identity, node and peer interchangeably.

3.2 Anonymous P2P Model

The system under attack is modelled as a general distributed computing environment where there is no constraint




Figure 4: The model in many sybil defence mechanisms can be seen as a social graph that is partitioned into a sybil region and an honest region. The two regions are connected by *attack edges*. Note that in general there may be multiple honest regions multiple sybil region.

on the topology, every node has limited computational resources and messages are guaranteed to be delivered [20]. Defending against the sybil attack is difficult under this model because nodes can communicate with any other node without authentication or authorisation.

3.3 Social Network Model

The more common model, especially in the context of social networks, is shown in Figure 4. It is first introduced by the authors of SybilGuard [90]. Nodes inside the left region are identities created by honest entities, the edges connecting those nodes are real-world trust relationships. The right region contains the sybils and they are connected with fake relationships. The edges connecting the two regions are called *attack edges*. These can be created by tricking an honest user to befriend a sybil, stealing an honest user’s account and so on. We call the nodes in the honest region that has one or more attack edges *gullible nodes* or *victim nodes*. An important property of this model is that the number of attack edges is not proportional to the number of sybils, but they are proportional to the number of honest nodes. So if attackers create too many sybils, then the graph will begin to have a quotient cut, that is removing the attack edges will disconnect a large number of nodes from the social graph. Many sybil defence mechanisms rely on the fact that attack edges are difficult to create as we will describe in section 5.

3.4 Impossibility Results

Under the anonymous P2P model, Douceur show that preventing the sybil attack is a lot more difficult because P2P systems often do not have a central, trusted authority. In fact, the author prove that is it impossible to completely prevent the sybil attack in a pure P2P environment [20].

Cheng and Friedman proved two results in the context of reputation systems [11]. Reputation systems are commonly used in e-commerce websites and the internet in general, where identities are rewarded by their good behaviour or usefulness. Google’s PageRank [59] is an example of a reputation system, where many links to a website makes it more reputable. The first result is that symmetric reputation

systems—that is when the reputation score does not depend on the names of the identities (this preserves anonymity)—cannot be made to prevent the sybil attack. The second result is for asymmetric reputation systems—that is when the reputation score for a node may be different depending on which node computed it—can be made to be sybil-resistant.

4. DERIVATIVES OF THE SYBIL ATTACK

The sybil attack can be seen as an umbrella term for all attacks that require the use of sybils. This section categorises the derivatives of the sybil attack by the attacker’s aim. Some attacks are more general, such as spamming, others are application specific, i.e. astroturfing. The attacks are sorted by their generality in decreasing order. We hope this section further illuminates the alarming consequences of the sybil attack.

4.1 Distributed Denial of Service

The DDoS (distributed denial of service) attack is where the attacker uses multiple identities to overload the target and deny it from doing work for its intended users.

In the BitTorrent network, DDoS attack can be directed at any machine connected to the internet, not just machines in the network. It uses a vulnerability in the network and a lot of sybils. The main idea is to report the victim as the tracker (a server that coordinates the peers) such that peers that are looking for a particular file would all query the victim [71].

El Defrawy et al. created a small scale proof-of-concept attack. Using only one machine, they could generate enough traffic to cripple small organisations and home users. The authors suggested that if sybils are created to perform the same attack aimed at a single victim, then it could easily throttle links with much higher bandwidth [21].

Steiner et al. also succeeded in mounting a DDoS attack but in the context of a file sharing DHT [76]. Instead of replying the correct list of peers to DHT queries, the sybils always respond with the IP address of the target peer in an attempt to overwhelm the target. The authors show evidence that real-world malicious DDoS attacks involving more than 300,000 peers are mounted using P2P networks.

4.2 Spamming

Spamming is the act of sending unsolicited or undesired messages (spam). The goal of the attacker varies from advertisement to phishing and spreading malicious software [32, 78]. Much of the spam originate from sybils because the attacker does not want to reveal their main identity and more importantly to circumvent blacklists. In the context of email, spam can be prevented in most cases when a large service provide is involved, such as Gmail [30]. The same is not true in other applications.

For social networks, many studies have characterised the behaviour of the spammers and found that many of them are automated sybils [77, 87, 28, 36]. More importantly, spamming is possibly the most common attack. Jiang et al. [36] analysed the malicious activities of the sybils on Renren⁴ and found that propagating advertisement is the most common activity (at 32.8%). Some authors have worked with the service provide to close the spam accounts, but it is

⁴One of the largest social network in China (<http://renren.com/>).

clearly not sufficient as we described in section 2.

In a very different application—cryptocurrencies, the spammer’s goal is to waste system resources rather than advertisements. Block sizes in Bitcoin is fixed to 1 MB and block generation takes about 10 minutes. Thus, there is an upper bound for the number of transactions that the Bitcoin network can handle per second (approximately 7 at the time of writing [5]). Flooding the network with useless transactions will cause its performance to drop. Decreasing the scalability even more. Fortunately, Bitcoin incorporates a transaction fee, so spammers cannot abuse the network indefinitely [6].

4.3 Slandering

The goal of a slandering attack is to illegitimately produce negative feedback to undermine the reputation of the target. It is mostly found in reputation systems where users are rewarded for their positive behaviour but also penalised for their bad behaviour. For example, if a buyer does not pay the merchant promptly in an online marketplace then the merchant may give a negative feedback to the buyer. The attack is difficult to prevent when the feedback is not authenticated. Furthermore, an attacker who uses a group of sybils can make the attack much more effective [33].

Mike Hearn suggests that Google has observed the attack in AdWords. The attackers would submit false reports on AdWords advertisements of their competitors to give them a bad reputation and gain an advantage over them [30].

Slandering also exist in wireless ad-hoc networks. Nodes in the network need to exchange information with each other to satisfy the underlying requirements of the application. Some common tasks are data aggregation and voting. With enough sybils, it is possible to manipulate the aggregated data or the poll to benefit the attacker. For example, sensor networks may use a ballot to detect misbehaving nodes, the attack could use its sybils to claim that an honest node is misbehaving and have the other nodes expel it from the network [58].

4.4 Self-promoting

In self-promotion, the goal of the attacker is to illegitimately raise its own reputation. It is also found in reputation systems, but it does not need to have penalisation property (e.g. the number of likes on a Facebook page rather than a decimal rating for a merchant between 1 and 5). A common way to perform self-promotion is to create sybils and have them create positive reputation for the attacker’s main identity.

De Cristofaro et al. performed an empirical study on Facebook page promotion using like-farms [15]. Some farms such as *SocialFormulae.com* are clearly operated by bots and the operator does not attempt to hide it, others such as *BoostLikes.com* tries to mimic human users. The authors purchased the “1000 likes” service on their empty Facebook pages. In under a month, many empty pages have accumulated almost 1000 likes as promised by the like-farms. A month later, the author’s empty accounts were not terminated, only a few of the liker’s account were terminated.

SEOClerks and MyCheapJobs are also evidences of marketplaces for self-promotion. Farooqi et al. found some of the top services include “1 million Twitter followers” at \$849, “1000+ Instagram followers” at \$10 and so on. The revenues of those two marketplaces are estimated to be at \$1.3 million

and \$116 thousand, respectively [23]. Although the authors did not investigate the properties of the fake followers, there is little doubt that many of accounts used in these services are sybils.

In the same vein as the slandering attack, wireless ad-hoc networks also suffer from the self-promoting attack. Nodes in the network often have limited resources such as bandwidth of the radio channels. Resources such as these must be shared between the neighbours using time slices. When the neighbours are sybils, then the attacker can receive an unfair amount of resource allocation, giving it a much higher privilege [58]. This works even when the sybils are themselves following the protocol.

4.5 Whitewashing

Whitewashing is when attackers abuse the reputation system for temporary gain and then escape the consequences by joining the reputation system under a new identity to shed their bad reputation [49]. For example, a Reddit user may delete his or her account when it accumulates a lot of negative karma, and then create a new account. The attack can be performed very easily if it is cheap to create sybils [33].

Whitewashing is a problem in file sharing networks that involve a central component. To prevent and penalise selfish behaviour such as free-riding, the central component acts as a reputation system. But many studies have found that the reputation system in fact encourages whitewashing [25, 88].

4.6 Routing Disruption

The aim of the attacker in routing disruption is to stop honest nodes from sending or receiving the correct data by maliciously modifying its intermediate routes.

In file sharing DHTs, routing disruption is commonly accomplished by index poisoning. That is when the attacker corrupts the routing tables so that honest peers fail to find the values they want. It can be mounted by injecting sybils into the DHT that do not follow the protocol. Wang and Kangasharju created honeypots in the BitTorrent network and detected that as many as 300,000 sybils may be performing such an attack [84]. Similar attacks are possible in other P2P networks such as Overnet [46].

In wireless ad-hoc networks, an important routing technique is multipath routing, where data is routed using multiple paths in the network for better fault-tolerance and bandwidth. However, if sybils are present in the network, then the different paths may in fact go through the sybils owned by a single attacker, nullifying the advantages. Another technique is geographic routing, nodes route data depending on the geographic location of their neighbours. Sybils in the network can pretend to be in more than one place at a time, thus significant disrupting the routing algorithm [39].

4.7 Eclipse Attack

The eclipse attack is closely related to routing disruption, sybils in this case are arranged in the network such that they *eclipse* the victim from the rest of the network. In other words, all of the victim's interactions are with the sybils. The victim can either be an identity or an object such as a key in a DHT [72].

Steiner et al. mounted an Eclipse attack on the KAD DHT, it is used in P2P file sharing networks Overnet and eMule [76]. The sybils are crafted such that they automatically position themselves as the neighbours of the target.

The authors show it is possible to eclipse a key using only 32 sybils in the DHT with 1.5 million users and 42,000 key.

Attackers can eclipse Bitcoin nodes to gain various unfair advantages [31]. Bitcoin has no authentication, so nodes receive unsolicited messages about the IP addresses of other nodes, these addresses are stored in an internal table. The attacker essentially sends bogus messages to the target that only contains the IP address of the attacker's sybils, or "rubbish" addresses. This technique eclipses the target because it can only connect to the attacker's IP address. This causes a lot of consequences. (1) Engineering block races, when two nodes discover the next block at the same time, the eclipsed block will not be able to receive the reward. (2) Eclipsing a large portion of the network can cause a 51% attack, where a single party gain complete control of the network. (3) Selfish-mining, that is when the attacker do not publish the latest block immediately after discovery but aims to publish 2 or more blocks at the same time to gain a lead; this results in many orphaned blocks for the other miners and the attacker can use the eclipse attack to drop newly discovered blocks by other miners to make the attack more effective. (4) Double spending, the eclipse attack allows the attacker to double spend his Bitcoins to eclipse nodes because they don't have an accurate view of the whole network.

4.8 Distributed Eavesdropping

Eavesdropping is the act of secretly recording the activities of one or more nodes. It is made harder by the nature of distributed systems, because there is no longer a central point of control. But sybils can be useful to overcome this problem.

Many authors have used sybils to monitor a P2P file sharing network that uses DHT [34, 76], such techniques can also be directly applied to eavesdrop users. In essence, the authors created a lot of light-weight sybils and tricked all the honest peers to store their addresses in their routing table, a form of index poisoning. The sybils are light-weight because they do not follow the DHT protocol and perform much simpler operations. A single machine can have thousands of sybils running simultaneously. Finally, DHT requests would "traverse through" the sybils due to the poisoned routing table, and the requests are stored in a database for further analysis.

4.9 Automated Identity Theft

Identity theft is the intentional use of the victim's identity to gain an advantage. Authors of [4] created two attacks - profile cloning and cross-site profile cloning, targeting five social network sites including Facebook and LinkedIn. The iCloner system was created to automate these attacks.

In profile cloning, iCloner uses publicly available information to automatically create clones of the victim's profiles. iCloner then sends friend requests from the cloned profile to the friends of the victim. The fact that the victim may have many friends that they do not contact very often, e.g. friend from primary school living in another country, makes this attack highly effective. The authors found that the acceptance rate for cloned profiles was over 60%. Much higher than the acceptance rate of 30% for fictitious profiles. Once the friendship is established, it is possible to extract private information that is not available publicly and perform identity theft.

The idea of cross-site profile cloning is similar, except the

cloned profile is created on another social network site that the victim does not yet use. Once the profile is cloned, iCloner attempts to identify friends of the victim and begins sending friend requests. Similarly, 56% of the friend requests were accepted.

More recently, Boshmaf et al. created SbN (Socialbot Network), which targets Facebook [9]. Each socialbot is a sybil created by the attacker, it controls a forged profile and mimics human behaviour to avoid detection. The attacker is the botmaster who coordinates the socialbots to achieve a common objective such as infiltrating the target social network by creating friend relationships with real users. The authors found that infiltration success rate was as high as 80% and the FIS (Facebook Immune System [75]) was not sufficient to prevent the attack. Once the relationships are established, the botmaster can command the socialbots to start gathering private information which can then be used for identity theft.

These examples demonstrate that the carelessness of users and the ability to create sybils makes social networks vulnerable to identity theft. Moreover, identity theft is only an entry point. Once trust relationships are established, the attacker can perform many other types of attacks such as spamming, phishing or astroturfing to gain advantage.

4.10 Astroturfing

Astroturfing is an act of creating grassroots movement that are in reality carried out by a single entity, effectively spreading misinformation to legitimate users. It relies on the ability to create sybils in the underlying social network. The 2010 Massachusetts senate race mentioned in the introduction is an example of such an attack.

Shortly after, Ratkiewicz et al. published their work on the Truthy system [63]—a web service that perform real-time analysis of Twitter to detect political astroturfing in the 2010 U.S. midterm election. The authors found accounts which generated a lot of retweets but no original tweets. More importantly, they uncovered a network of bot accounts that injected thousands of tweets to smear the Democratic candidate.

In 2012, Wang et al. investigated two of the largest crowd-turfing (crowdsourced astroturfing) platforms in China—Zhubajie and Sandaha. One of their services is to perform astroturfing on Weibo⁵. The authors found that the 5364 sellers collectively own 14151 Weibo accounts and the top 1% of the sellers own over 100 accounts. Furthermore, the business is growing and more than \$4 million have been spent on these two platforms over five years [83].

5. SYBIL DEFENCE MECHANISMS

The sybil attack clearly has a lot of consequences, some of them are severe and may cause financial losses or undermine freedom of speech. In this section we categorise various defence techniques against the sybil attack. We classify them on their main idea, and state explicitly when the mechanism is application specific. An overview sorted by the year of first publication is provided on Table 1. The articles are selected based on their impact (e.g. the number of citations and the usefulness in the real world) and novelty. We attempt to include all the seminal work.

⁵The Chinese Twitter (<http://weibo.com/>).

5.1 Reputation Systems

Some reputation systems exhibit the ability to resist sybils. The canonical reputation system is PageRank [59]. In the context of world wide web, it assigns a score to every web page depending on the number of links pointing to it from other web pages. Suppose an attacker wants to boost the score of his web page, he would create sybil pages and create links to his main page. PageRank prevent this type of manipulation because the sybil pages do not have a high score so they cannot influence the attacker’s page by a large amount [2]. On the other hand, PageRank is not immune to the sybil attack because the initial scores for the sybil pages are not zero. Cheng and Friedman found that for about 300,000 pages it requires 500 sybils to rise a median node to the top 100 [12].

Many other reputation systems exist other than PageRank [38, 74, 57]. However, many of them do not factor the sybil attack into their design and suggest using an independent defence mechanism which we describe in the following sections.

5.2 Certificate Authority

CA (certificate authorities) check the users’ real identities and then issues certificates to honest users. The certificate can be tangible (trusted hardware [58]) or non-tangible (public key certificate) depending on the application. When an identity wishes to use the application, the CA must verify the validity of its certificate to ensure one-to-one correspondence. This mechanism prevents the sybil attack as long as the CA does not make mistakes in the issuance stage.

Many existing systems today use a form of CA. X.509 [35]—a standard for certificates, it is used in a large variety of applications. For example, emails can be encrypted and signed using S/MIME [62] certificates which are based on X.509. Attackers can still create many sybils and send emails, but the receiver would reject the emails because they are not correctly signed.

CA can prevent the sybil attack but it also has a lot of downsides. (1) Users have different opinions and may not agree on a single CA. (2) Users living in authoritarian regimes may not have access to the necessary CA. (3) It is difficult to scale up a CA to meet increasing users demands. (4) Anonymity is difficult to obtain because the CA has complete information of the entities. (5) It is a central point of failure; i.e. if the attacker obtains the private key to create certificates then he or she can easily generate sybils, if the CA goes offline then the application ceases to function because it can no longer verify identities.

5.3 Resource Testing

Resource testing makes attacks costly. That is to say every attacker can create multiple sybils, but the attack cannot duplicate its resources the same way. The resource type varies depending on the application, and we give a few examples below. It may deter casual attackers but its usefulness degrades for resourceful attackers.

5.3.1 IP Addresses

In P2P networks, IP address can be used as a resource. In Tarzan [26], neighbours are selected not from all known IP addresses, but from distinct IP prefixes. The effectiveness of the sybil attack is reduced if the attackers cannot easily create sybils in a large range of IP prefixes. Another

Table 1: Overview of the defence mechanisms surveyed in this work.

Year	Mechanism	Type			Classification									
		Restrict sybil creation	Perform sybil detection	Restrict sybil participation	Reputation system	Certificate authority	Resource testing	Registration fee	Network flow	Social networks	Community detection	Content based	Machine learning	Miscellaneous
1999	PageRank [59]			•	•									
2001	Once-in-a-lifetime Identity [64]	•						•						
2002	X.509 [35]	•				•								
2002	Test IP Address [26]	•					•							
2003	CAPCHA [82]	•						•						
2004	Trusted Hardware [58]	•				•								
2004	Adaptive Stranger [24]			•				•						
2005	Test CPU - Puzzle [1]			•			•							
2005	Trust Transfer [66]			•										•
2006	SybilGuard [90]		•											
2008	Test CPU - Bitcoin [56]			•			•							
2008	SybilLimit [89]		•							•				
2008	Ostra [51]			•								•		
2009	BarterCast [50]			•					•					
2009	SumUp [79]			•					•					
2009	SybilInfer [14]		•							•				
2009	DSybil [91]			•										•
2009	SyMon [37]		•											•
2010	Whānau [45]			•						•				
2010	Mislove’s Algorithm [81]		•								•			
2010	Stringhini et el. [77]		•										•	
2011	DropEdge [68]			•					•					
2011	Sparse Cut [42]		•						•					
2011	GateKeeper [80]			•					•	•				
2012	SybilRes [16]			•					•					
2012	SybilDefender [85]		•							•				
2012	SybilRank [10]		•							•				
2013	SybilShield [70]		•							•				
2013	VoteTrust [86]			•								•		
2015	Integro [8]		•							•			•	
2016	SybilExposer [53]		•								•			
2016	SmartWalk [48]		•							•			•	

example the self-registration technique [17]. When a peer wish to join the network, it needs to compute an ID which is a cryptographically secure hash of its own IP address and port number, and then broadcast it to the network. While participating in the network, other peers need to verify that the ID matches the peer's origin.

5.3.2 CPU

Aspnes et al. proposed the idea of solving difficult puzzles to limiting the number of sybils [1]. The puzzle in this case is computing hashes on some input y concatenated by a some string x such that the digest begins with w number of zeros. Essentially, every node acts as a verifier by picking y and then broadcast it to the network to the puzzle solver⁶. The puzzle solver must compute as many x as possible such that they match the requirement. Sybils are unlikely to produce enough x 's so the honest nodes will refuse to interact with them.

In fact, many crypto-currencies use the same idea. For instance *proof-of-work* in Bitcoin [56]. The Bitcoin blockchain is a global ledger and it needs to reach a consensus for its blocks. Nodes in the Bitcoin network essentially "vote" for the latest block. But the vote is performed not by counting the majority, but by "counting" the amount of CPU power, i.e. one CPU is one vote. Thus, an attacker cannot simply create a lot of identities to out-vote the honest nodes. It needs to gather a lot of CPU power which is a lot more difficult.

5.4 Registration Fee

Friedman and Resnick is one of the first to propose the use of a registration fee [64]. It is similar to resource testing except it only happens on registration. Entities can be charged a fee for creating identities, often facilitated by a central authority. The fee needs to be set appropriately so that the cost of creating sybils outweighs the benefits but does not hinder honest entities.

5.4.1 Once-in-a-lifetime Identity

The fee does not need to be monetary. Friedman and Resnick proposed the idea of a once-in-a-lifetime identity [64]. It uses blind signatures and a central authority, the authority does not know the mapping between the real identity (e.g. driving licence) and the pseudo identity, but it checks whether there has been previous registrations of the same real identity. In this case, the fee is the real identity, and attackers cannot create an arbitrary number of real identities.

5.4.2 CAPTCHA

CAPTCHA [82] prevents programs from automatically creating new identities and limits the rate at which identities can be created by asking users to solve a puzzle that is difficult for computers. In this case the registration fee is the time required to solve the puzzle.

5.4.3 Adaptive Stranger Policy

Feldman et al. proposed another form of registration fee for P2P networks - the adaptive stranger policy [24]. When new peers join the network, they are treated using a policy that is adapted from previous newcomers. For example, the

⁶We simplify the protocol a bit because the original protocol (called Democracy) is made to be Byzantine fault tolerant and is a bit more involved.




Figure 5: Subjective graph of A. The numbers are the amount of data transferred, they can be seen as the capacity in the context of the maximum flow problem.

new peers may be expected to contribute to the network before they are allowed to receive benefits from the "mature" peers. The downside is that the policy may deter honest users from joining the network in the first place. The fee in this case is the initial contribution.

5.5 Network Flow

Network flow based techniques began with BarterCast [50]. It was initially designed to combat freeriding in P2P file sharing networks, where users are selfish and do not share content, but its idea can be extended combat the sybil attack. The ideas based on BarterCast do not directly identify sybils, but they prevent sybils from doing harm in the P2P network, we refer these are sybil tolerance systems.

5.5.1 BarterCast

The main idea of BarterCast comes from human interactions, where the reputation of a person can be from direct experiences, or information obtained from someone else. As we know, our direct experiences are always true, but the indirect information may not be, i.e. people can lie about their experiences. Humans solve the problem by treating the indirect information with a grain of salt unless the source of the information is highly trusted.

BarterCast applies this idea in P2P file sharing networks. Peers all maintain a subjective graph which is created by exchanging messages with their neighbours. The direct experiences, measured by the number of bytes uploaded and downloaded are represented by the outgoing and incoming edges from the peer, respectively. Indirect experiences are represented by edges that are not directly connected to the peer. For example in Figure 5, A is the subject, it has direct experiences with B and B has told A about S , so it has indirect information about S . But A is unsure about the truthfulness of S 's contribution, so it only trusts S as much as it trusts B . This idea is realised using the Ford-Fulkerson maximum flow algorithm [13] as shown in Figure 5, A only have 5 units of trust for S even when S apparently contributed a lot to B .

BarterCast does not prevent the sybil attack by itself. Attackers can first upload a lot of data to obtain a good reputation in the network, and if the attacker creates sybils and lie to the honest node that the sybils contributed a lot. Then the peers who have interacted with the attacker will be tricked to think that the sybils also have a high reputation.

5.5.2 SybilRes

To fix the problem of BarterCast, Delaviz et al. created SybilRes [16]. The main idea is the following. Suppose there are two peers A and B who are sharing data. If A is uploading (represented by an outgoing edge) to B , then it decreases the weight of the incoming edge from B . Vice versa, the weight is increased for the outgoing edge when A is downloading. The rate of change depends on the capacities of the edges and the amount of data transferred after computing the reputation. Using the definition in Figure 4, the attacker cannot build up reputation for its sybils by uploading to peers in the honest region beforehand, it is now forced to keep on uploading to keep its sybil’s reputation which is a more desirable behaviour.

5.5.3 DropEdge

Seuken et al. provided a formal model of BarterCast. They found that BarterCast is vulnerable to misreporting and proposed a solution called the DropEdge mechanism [68, 69]. DropEdge, like the name implies, drops some edges in the subjective graph that satisfies the following constraints. Suppose peer A wishes to download from peers in set C (the choice set). Then any reports received by A from $p \in C$ is dropped. Also, edges with both end points in C are also dropped from A ’s subjective graph. Peers in C cannot misreport their contribution since all the necessary edges are dropped. The authors prove that it is robust against weakly beneficial sybils, that is sybils that do not perform actual work for honest peers. The authors also prove that no mechanism can prevent strongly beneficial sybils, i.e. sybils that interact with honest users.

5.5.4 SumUp

SumUp [79] is a defence mechanism specific for the vote aggregation problem. For example, in social news aggregation websites such as Reddit, users vote on the submitted content to determine its ranking; the problem occurs when sybils can out-vote honest users. It is a centralised approach that fits the architecture of most websites that perform vote aggregation. SumUp consist of three stages. Firstly, pruning is performed to limit the number of incoming edges of every node, this is to reduce the number of attack edges available and reduce the computational cost in later stages, the threshold for the number of edges to prune is a system parameter. Secondly, it uses a ticket source (the central component) to distribute tickets in a breadth-first search manner equally to its neighbours, every node keeps one ticket and distributes the remaining tickets the same way. The number of tickets distributed across an edge plus one is the capacity of the edge. Effectively, edges closer to the ticket source have a high capacity. This idea keeps the capacities in the sybil region low so that they do not have a large influence on the outcome. Finally, the maximum flow is computed where the source is simply the ticket source and the sink is an imaginary node with edges of capacity one that is connected to every voter. SumUp offers a better guarantee than SybilLimit (one of the primary defence mechanisms in social networks which we describe in subsection 5.6) where it only accepts $1 + o(n)$ votes per attack edge. Unlike the aforementioned techniques in this section, SumUp requires a social network so it does not work in a generic P2P setting. An improved version of SumUp - GateKeeper is discussed in subsection 5.6.




Figure 6: Visualisation of the ticket distribution step of SumUp. Note that tickets are not distributed to nodes on the same level of breadth-first search, or nodes that already hold a ticket. The same process is also applied for GateKeeper, discussed in at the end of subsection 5.6.

5.5.5 Sparse Cut

Conversely, maximum flow is dual to minimum cut, so the problem of finding sybil can also be formulated as finding sparse cuts. The sparse cut problem is to find a partition such that the ratio between the number in the cut and the number of vertices in the smaller partition is minimised. Kurve and Kesidis devised an algorithm for finding sparse cuts to detect sybils [42]. It relies on the presence of trusted nodes.

5.6 Random Walk in Social Networks

Another family, possibly the largest, sybil defence mechanism is based on random walks in social networks, first proposed in SybilGuard [90]. The key assumptions in these techniques is that the honest region is *fast mixing*⁷, and the attack edges are difficult to form and are independent of the number of sybils.

Before explaining the techniques, we define the terms *random walk* and *random route* in the context of social graphs. In random walk, the social graph is traversed such that outgoing edges are selected uniformly at random on every hop of the walk. Random route is a modified form of a random walk. Every node maintains a static routing table that contains a uniformly random one-to-one mapping between incoming edges and outgoing edges, initialised at start-up. Thus, a route is determined by the tables on every node. An important property of random route is that if two routes enter the same edge, then they will always exit at the same edge, so their route after exiting will be exactly the same. In most cases, the number of hops for a random route should be just right, so that the fast mixing property is achieved in the honest region, this is known as the *mixing time*.

⁷In a graph, if a random walk of length $O(\log N)$ reaches a stationary distribution of nodes, then the graph is fast mixing.




Figure 7: Visualisation of SybilGuard. The verifier accepts the suspect because their random routes intersect. The Sybils’ random routes all come from a single attack edge, thus their routes are equivalent after entering the honest region.

5.6.1 SybilGuard

In SybilGuard [90] (visualised in Figure 7), every node acts as a verifier and performs a single random route of a fixed length, determined by the mixing time. The verifier treats every other node as suspects initially. The suspect is labelled as an honest node if its random route intersects with the verifier’s random route. The number of accepted nodes for every intersection is limited by a quota. Intuitively, the random route from an honest node is unlikely to escape into the sybil region because the number of attack edges is limited. The number of overlapping random routes from the sybils is bounded by the number of attack edges due to the random route property. Recall that the number of attack edges is independent of the number of sybils, thus they are unlikely to intersect with many honest nodes.

5.6.2 SybilLimit

SybilLimit [89] is the continuation of SybilGuard and it is an improvement on many fronts while keeping the same or better guarantees. Same as before, every honest node acts as a verifier V and initially treats all other nodes as suspects S . The verification process begins by performing multiple independent random routes instead of a single one as in SybilGuard. V labels S as an honest node if and only if they share at least one tail (the final edge in the route). For each tail of V , there is a quota for the number of node that it labels. The authors prove that SybilLimit bounds the number of accepted sybils (false positives) at $O(\log n)$, an improvement from $O(\sqrt{n} \log n)$ of SybilGuard. The process is visualised in Figure 8.

Let us consider the following three scenarios to intuitively show why SybilLimit works. The same intuition applies to SybilGuard. Suppose S is not a sybil, and if V and S perform enough random routes, each with enough hops for fast mixing, then due to the Birthday Paradox, S and V will have an intersecting tail with high probability. Next, suppose some tails of V are in the sybil region so they may intersect with many sybils, but crossing the attack edge is improbable and accepting a lot of sybils is also difficult due to the aforementioned quota mechanism, thus V has a small probability of accepting a large number of sybils. Finally,




Figure 8: Visualisation of SybilLimit. The verifier accepts the suspect because one of the tails of their random routes meet. Similarly, the Sybils’ random route are all from a single attack edges, so their routes are equivalent in the honest region.

consider there is only one attack edge and suppose a sybil has tails in the honest region, due to the random route property, the route of the sybils in the honest region will be equivalent (overlapping), so accepting the sybils in this scenario is also low due to the quota mechanism.

5.6.3 SybilInfer

SybilInfer [14] is inspired by SybilGuard and SybilLimit. It assumes trusted nodes, which create traces by doing random walks in the graph. Based on the traces, a probability model that describes the likelihood a trace T was generated by a specific set of honest nodes X , i.e. $\Pr[T|X = \text{honest}]$. Then using Bayesian inference, $\Pr[X = \text{honest}|T]$ can be computed, that is effectively assigning a “score” to every node. Sybils are the nodes with a low “score”. SybilInfer outperforms SybilLimit regarding the number of false positives, but its drawbacks are its high computational cost and reliance on trusted nodes.

5.6.4 Whānau

Lesniewski et al. created Whānau, a sybil-resistant DHT, it is also inspired by SybilLimit [44, 45]. Suppose all nodes in the DHT belong to a social network and a node, say Alice (who is honest), wish to join the DHT. Alice performs multiple random walks and it inserts the node on the tails of her random walks into her own routing table. Again, the random walk properties guarantee that there cannot be a large proportion of sybils in the routing table with a high probability. The number of entries in the routing table needs to be high enough such that the entries cover the whole key spaces. If Alice wants to find a key, she broadcasts the request to all the entries. Since most of the entries are honest, Alice can retrieve the required value with high probability.

5.6.5 SybilDefender

SybilDefender [85] can be seen as a two step process. It assumes the size of the sybil region is smaller than the honest region and the nodes in the sybil region are well connected. The first step is to perform random walk to detect sybils, similar to SybilGuard. The second step is to detect a complete sybil region around the detected sybils. It performs a

partial random walk, where the random walk is not allowed to traverse the same node more than once. A property of partial random walks is that they are likely to “die” (all the neighbour nodes have already been traversed) upon reaching the edge of the sybil region, thus they are likely to stay in the sybil region. The sybil region is detected by examining the nodes traversed by the partial random walk.

5.6.6 SybilRank

SybilRank [10], in contrast to the aforementioned techniques, is designed to be integrated with real-world social networks, and is deployed on Tuenti⁸. SybilRank uses short random walks that begins on trusted nodes in the honest region. The trusted nodes is chosen manually, this allows SybilRank to adapt to different graph structures. A novelty in SybilRank is that it uses power iterations (like in PageRank) for computing the landing probability of random walks. Intuitively, the landing probability decreases for nodes that are far away from the trusted nodes (since it is using short random walks), especially for nodes in the sybil region. The probabilities are normalised by the degree of the node and then ranked. The potential sybils are the nodes that are under some threshold—a system parameter. Finally, various actions can be performed to verify the potential sybils, e.g. using CAPTCHA puzzles.

5.6.7 SybilShield

SybilShield [70] makes use of multiple communities (Figure 9). It begins the same way as SybilGuard/SybilLimit, i.e. V performs random route to determine whether suspect S is a sybil. But to reduce the possibility that S is in fact an honest node but labelled as a sybil, V searches for agents A that are from another community. This is also done using random routes and relies on the assumption that inter-community edges are rare. To do this, V performs a random route and picks a *candidate* A , then V and the candidate perform random routes simultaneously, if they do not intersect then A is considered to be in another community, otherwise V repeats the process until it finds a suitable A . When a number of suitable A is found, they all perform random route and decides whether S is actually a sybil and then relay the information back to V . If a large majority of A say S is honest, then V knows that it has made a mistake, otherwise S is indeed a sybil.

5.6.8 SmartWalk

All the techniques so far use fixed length random walks. Recently, Liu et al. argue that fixed length random walks is not adequate because social graphs can have communities of different sizes, nodes in a large community that is well connected may have a lower mixing time than nodes in a smaller community. A longer than desired random walk length will result in a growing number of false negatives. To solve this problem, the authors created SmartWalk [48] which uses adaptive random walks. There are two steps to SmartWalk. The first step uses machine learning to predict the local mixing time given a node. The second step performs the random walk using the local mixing time as the starting length, but on every hop the number of hops remaining is also updated depending on the intermediate node’s local mixing time. Finally, the sybils can be detected using the same method as SybilLimit, i.e. node is accepted

⁸A Spanish social network with 11 million users




Figure 9: Visualisation of communities in a social graph. Solid lines represent intra-community edges. Dotted lines represent inter-community edges. The red dotted lines are inter-community edges coming from the sybils, in other words the attack edges. The idea of making use of communities is used in SybilShield as well as SybilExposer (subsection 5.7).

when their tails meet. The authors evaluated SmartWalk using real-world social graphs. In particular, the false positive rate for a Twitter social graph can be reduced by two orders of magnitude when compared to the standard SybilLimit.

5.6.9 GateKeeper

GateKeeper [80] combines ideas from SumUp (discussed in subsection 5.5) and SybilLimit. GateKeeper assumes an admission controller that is honest, the admission controller performs random walks to select n ticket sources. The ticket sources act the same way as SumUp where it distributes ticket in a breadth-first search manner (Figure 6). For a node to be labelled honest, it must obtain fn tickets, where f is a system parameter (0.2 is shown to be a good value experimentally). This idea works because if the ticket sources are evenly distributed and sybils only have a few attack edges. Thus it is unlikely that the sybils will receive many tickets.

5.7 Community Detection

In this section we discuss techniques that leverage existing community detection algorithm. The work started with Viswanath et al., who realised many of the mechanisms mentioned in subsection 5.6 such as SybilGuard, SybilLimit and SybilInfer, are in fact performing local community detection (i.e. detecting clusters of nodes) which is a more developed field [81]. The authors also argue that social graphs are not always fast mixing, which may result in poor results for techniques that use the fast mixing assumption.

5.7.1 Mislove’s Algorithm

Viswanath et al. applied the Mislove’s community detection algorithm [52] on synthetic social graphs to detect sybils [81]. The results show it is just as effective as SybilLimit and SybilInfer, but it is much better than SybilGuard. But using real data, namely a Facebook social graph, Mislove’s algorithm performs better than the other methods.

The result of SybilGuard in this case is nearly identical to picking sybils at random.

5.7.2 SybilExposer

The authors of SybilExposer [53] argue that the number of attack edges may not be that small which may render the random walk based method ineffective. They proposed a solution that relies on the ratio between the number of inter-community (inter-cluster) edges and the number of intra-community (intra-cluster) edges. This is visualised in Figure 9 where the sybil region has fewer inter-community edges than the honest region. The idea is that this ratio is different between honest communities and sybil communities, namely the sybil communities have a lower ratio because they are well connected between themselves but not with other honest communities. SybilExposer operates in two stages, first communities are extracted using community detection algorithm (a modified version of the Louvain method [7]), then the communities are ranked based on the ratio and communities with a low ratio are likely to be sybils.

5.8 Content Based and Machine Learning

In some application domains such as social networks, it is possible to leverage user content or user feedback to detect sybils. These techniques work well in practice. But often depend on uninformed attackers that do not try to mimic the behaviour of honest nodes.

5.8.1 Ostra

Ostra [51] is a system for reducing spam in social networks. In the simplest form, every undirected edge in the social graph is considered as two directed edges, each of them has a credit value. When a user wants to send a message, it needs to find a path with enough credits in the social graph from itself to the receiver. The edge traversed by the message will have its credit deducted, and the opposite edge will have its credit added. The receiver then decides whether the message is a spam, if it's not a spam then the credit operations are reversed. Effectively, only spam messages will have an effect on the credits. If a path cannot be found, i.e. all possible paths have run out of credit, then the message is blocked. Naturally, spam from the sybils must use the attack edges, if enough honest users mark those messages as spam then the credit on the attack edges will run out and the sybils can no longer send messages.

5.8.2 Machine Learning on Honey-profiles

Stringhini et al. devised a machine learning technique to classify bot accounts in Twitter [77]. Six features were devised. One of them is "FF Ratio", that is the ratio between number of users that the account is following and the number of followers. Other features include "URL Ratio", "Message Similarity" and so on. The authors collected data on "honey-profiles", trained a classifier after analysing those data and collaborated with Twitter to delete tens of thousands of spam accounts.

5.8.3 VoteTrust

VoteTrust [86] leverages the distrust relationship, i.e. friend request rejection, to detect sybils in social networks. Suppose A sends a friend request to B , if B accepts/rejects the request then it is considered as a positive/negative vote on A by B . The first step is to use PageRank combined

with human scrutiny to select a number of trust seeds in the honest region. Then the trusted seeds distributes *vote capacity*, that is the number of votes each node can cast. Initially only the trusted seeds have a positive vote capacity and other nodes have 0. When a node receives a positive vote from a trusted seed, it also receives some vote capacity. Then it can repeat the same process on nodes it votes on, thus distributing the vote capacity. The vote capacity decreases as it goes further away from the trusted seeds. This technique is comparable to the ticket distribution technique used in SumUp and GateKeeper. Finally, the votes are aggregated for every node in the graph. Naturally, the sybils are likely to have a low vote because their vote capacity is low and many of their friend requests would be rejected.

5.8.4 Integro

Integro [8] is a hybrid between random walk and content based approaches. It begins by training a machine learning algorithm to identify potential *victim accounts*, that is honest accounts that have accepted sybils as their friends. Then in the social graph, edges connecting the potential victim accounts will have its weight reduced depending on the likelihood of it being a victim. Finally, it performs biased short random walk starting from some known honest account to compute the landing probabilities for every node. Biased in a sense that the walk is a higher probability of using a path with a higher weight. Sybils are the nodes with a low landing probability. This technique works because, victims are easier to detect than sybils due to the fact that sybils can arbitrarily modify their account information to avoid detection. Once the victims are detected, they effectively form a "border-line" between the honest region and the sybil region. Finally, it is unlikely that the random walk will traverse into the sybil region due to its bias, so the sybil will have a low landing probability and be detected.

5.9 Miscellaneous

There are many defence mechanisms that unique in their own right. We cover some of them in this section.

5.9.1 Trust Transfer

Trust transfer [66] is a sybil defence mechanism for reputation systems that transfers the reputation score from a recommender to a recommended identity. This method discourages self-recommendation behaviour because the attacker would need to lower the reputation of its sybils to recommend him or herself. The sybils cannot gain reputation from honest identities because if they do not interact with them. It may be strange to lose reputation when recommending an identity, but the authors argue that in certain scenarios where there are a lot of interactions and the overall trustworthiness is high, then there is no major effect to transfer a little reputation to a recommended identity.

5.9.2 DSybil

Yu et al. of DSybil [91] argue that defending sybils in reputation or recommendation systems is a lot more difficult than in social networks because only a very small percentage of the user will vote for an object (e.g. news article in Reddit), so a few sybils and attack edges can easily out-vote honest users. Their proposed solution is DSybil, a distributed algorithm for diminishing the influence of sybils in recommendation systems using historical data. Suppose Al-

ice is an identity that runs the algorithm, and every identity begins with the same trust score from Alice’s perspective. The algorithm runs in rounds. In every round, Alice picks an object to consume (e.g. reads the new article on Reddit) and then makes a binary (good or bad) feedback on the object. Then Alice computes whether the object is *overwhelming*, namely whether the sum of the trust scores of the voters of the object exceeds some threshold. If Alice voted for good and the object is not overwhelming, then she would increase the trust scores by some factor for all the voters of that object. Otherwise, she decreases the trust score by some factor. When Alice needs a recommendation, a uniformly random overwhelming object is returned. Trust scores for identities that have the same interest as Alice grow exponentially when Alice consumes a good non-overwhelming object. Conversely, the trust scores decreases exponentially for identities that are recommending bad objects, making sybils ineffective.

5.9.3 SyMon

SyMon [37] or *Sybil Monitor* assumes that any two sufficiently random nodes in the network cannot both be sybils with a high probability. Then the nodes are paired together to monitor each other’s transactions. For instance, a transaction could be reporting the number bytes transferred in a P2P file sharing network. Cheating occurs when a node reports some bytes transferred that does not match its network traffic. The authors provide four methods for pairing nodes. Suppose the nodes are identified by a cryptographically secure hash of their RSA public key, then it is difficult to create identities deterministically. Nodes can be matched by the closeness of their identities. The downside of this approach is that it sacrifices a lot of privacy, every action that a node makes is monitored by some other node.

6. SURVEY OF SYBIL ATTACK SURVEYS

Many surveys on the sybil attack exist in the literature. We attribute much of the initial findings to these surveys. We describe and compare them with our work in this section.

To the best of our knowledge, Levine et al. published the first survey of the defence mechanisms [49] in 2006. They found that the most popular defence mechanism (in terms of the number of published work) at that time is to use a certificate authority. The surveyed defence mechanisms approximately cover sections 5.2, 5.3 and 5.4 in this work.

Many more surveys were published after the introduction of social network and random walk based sybil defences beginning with SybilGuard. Mohaisen and Kim surveyed certificate authority, resource testing and social network based approaches [54]. They also compare the assumptions, performance and many other properties. Rakesh et al. made a similar survey [61], they discuss six types of attacks in addition to describing the defence mechanisms. Our work can be seen as an extended work of these surveys—in terms of the possible attacks, and the defence mechanisms.

Koll et al. surveyed 8 defence mechanisms for social networks and analysed them in much more depth [40]. The authors experimentally show that increasing the number of attack edges indeed makes the defence mechanisms less effective. More interestingly, sybil tolerance systems such as Ostra (discussed in subsection 5.8) and SumUp (discussed in subsection 5.5) can still be effective when the number of attack edges increase. The authors advise future defence

mechanism designers to use information in addition to simply the graph structure to detect or tolerate sybils. In comparison, our work do not go into the same level of depth, but provide a much broader spectrum of defence mechanisms.

Surveys of reputation often cover sybil attacks too [49, 33, 41, 67]. Although they do not cover the sybil attack in depth, these surveys provide a lot of insight from a different perspective, especially for the possible attacks in reputation systems.

7. CONCLUSION

In this work, we survey both the practical and the theoretical aspects of the sybil attack. We demonstrate the severity of the sybil attack by showing the harm it is causing in the real-world. Social networks are flooded with sybils making real news and fake news almost indistinguishable. Users of Tor are monitored by sybils hidden in the network trying to reveal their real identity. We illustrate the severity of the attack further by demonstrating the simplicity of purchasing thousands of fake Twitter followers. Next, we define the sybil attack using its original definition from Douceur [20] and introduce the models and definitions used in this work. With the necessary background, we describe the derivatives of the sybil attack, some alarming attack include self-promoting and identity theft. The main part of our work summarises the defence mechanisms. Earlier defence mechanisms primarily work by limiting the number of identities or the rate at which they are created. The introduction of BarterCast inspired many network-flow based techniques for limiting the influence of sybils. Similarly, the introduction of SybilGuard stimulated a lot of work on random-walk based techniques for identifying sybils in social networks. Hybrids are also available, for instance GateKeeper is a hybrid of the two aforementioned techniques. Finally, we compare and contrast our work with existing surveys.

We hope this work demonstrates the alarming consequences of the sybil attack and many ingenious ways to defend against it. However, there does not exist a general solution and many defence mechanisms must satisfy their own set of assumptions in order to perform well. When the assumptions are violated, which can be the case due to the dynamic structure of real networks, they become ineffective (demonstrated in [48]). Moreover, almost no defence mechanisms considers the temporal dynamics [47], that is when the attacker may modify the attack edges or its the social graph in the sybil region over time. Lin et al. show the attacker can “greatly undermine the security guarantees” of many defence mechanisms [47].

Without a doubt, much work still needs to be done in order for the cyberspace to be free of the sybil attack. We hope this work serves as a cornerstone for the future defence mechanisms.

8. REFERENCES

- [1] J. Aspnes, C. Jackson, and A. Krishnamurthy. Exposing computationally-challenged Byzantine impostors. *Department of Computer Science, Yale University, New Haven, CT, Tech. Rep.*, 2005.
- [2] R. A. Baeza-Yates, C. Castillo, and V. López. Pagerank Increase under Different Collusion Topologies. In *AIRWeb*, volume 5, pages 25–32, 2005.

- [3] R. Bhattacharjee and A. Goel. Avoiding ballot stuffing in ebay-like reputation systems. In *Proceedings of the 2005 ACM SIGCOMM workshop on Economics of peer-to-peer systems*, pages 133–137. ACM, 2005.
- [4] L. Bilge, T. Strufe, D. Balzarotti, and E. Kirda. All your contacts are belong to us: automated identity theft attacks on social networks. In *Proceedings of the 18th international conference on World wide web*, pages 551–560. ACM, 2009.
- [5] Bitcoinwiki. Scalability. <https://en.bitcoin.it/wiki/Scalability>, 2016. Accessed: 2016-11-23.
- [6] Bitcoinwiki. Spam transactions. https://en.bitcoin.it/wiki/Spam_transactions, 2016. Accessed: 2016-11-23.
- [7] V. D. Blondel, J.-L. Guillaume, R. Lambiotte, and E. Lefebvre. Fast unfolding of communities in large networks. *Journal of statistical mechanics: theory and experiment*, 2008(10):P10008, 2008.
- [8] Y. Boshmaf, D. Logothetis, G. Siganos, J. Lería, J. Lorenzo, M. Ripeanu, and K. Beznosov. Integro: Leveraging Victim Prediction for Robust Fake Account Detection in OSNs. In *NDSS*, volume 15, pages 8–11, 2015.
- [9] Y. Boshmaf, I. Musluhkhov, K. Beznosov, and M. Ripeanu. The socialbot network: when bots socialize for fame and money. In *Proceedings of the 27th Annual Computer Security Applications Conference*, pages 93–102. ACM, 2011.
- [10] Q. Cao, M. Sirivianos, X. Yang, and T. Pregueiro. Aiding the detection of fake accounts in large scale social online services. In *Presented as part of the 9th USENIX Symposium on Networked Systems Design and Implementation (NSDI 12)*, pages 197–210, 2012.
- [11] A. Cheng and E. Friedman. Sybilproof reputation mechanisms. In *Proceedings of the 2005 ACM SIGCOMM workshop on Economics of peer-to-peer systems*, pages 128–132. ACM, 2005.
- [12] A. Cheng and E. Friedman. Manipulability of PageRank under sybil strategies, 2006.
- [13] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein. *Introduction to algorithms*, volume 6. MIT press Cambridge, 2001.
- [14] G. Danezis and P. Mittal. SybilInfer: Detecting Sybil Nodes using Social Networks. In *NDSS*. San Diego, CA, 2009.
- [15] E. De Cristofaro, A. Friedman, G. Jourjon, M. A. Kaafar, and M. Z. Shafiq. Paying for likes?: Understanding facebook like fraud using honeypots. In *Proceedings of the 2014 Conference on Internet Measurement Conference*, pages 129–136. ACM, 2014.
- [16] R. Delaviz, N. Andrade, J. A. Pouwelse, and D. H. Epema. SybilRes: A sybil-resilient flow-based decentralized reputation mechanism. In *Distributed Computing Systems (ICDCS), 2012 IEEE 32nd International Conference on*, pages 203–213. IEEE, 2012.
- [17] J. Dinger and H. Hartenstein. Defending the sybil attack in p2p networks: Taxonomy, challenges, and a proposal for self-registration. In *First International Conference on Availability, Reliability and Security (ARES’06)*, pages 8–pp. IEEE, 2006.
- [18] R. Dingledine. Tor security advisory: “relay early” traffic confirmation attack. <https://blog.torproject.org/blog/tor-security-advisory-relay-early-traffic-confirmation-attack>, 2014. Accessed: 2016-11-16.
- [19] R. Dingledine, N. Mathewson, and P. Syverson. Tor: The second-generation onion router. Technical report, DTIC Document, 2004.
- [20] J. R. Douceur. The sybil attack. In *International Workshop on Peer-to-Peer Systems*, pages 251–260. Springer, 2002.
- [21] K. El Defrawy, M. Gjoka, and A. Markopoulou. BotTorrent: Misusing BitTorrent to Launch DDos Attacks. *SRUTI*, 7:1–6, 2007.
- [22] M. Faloutsos, P. Faloutsos, and C. Faloutsos. On power-law relationships of the internet topology. In *ACM SIGCOMM computer communication review*, volume 29, pages 251–262. ACM, 1999.
- [23] S. Farooqi, M. Ikram, G. Irfan, E. De Cristofaro, A. Friedman, G. Jourjon, M. A. Kaafar, M. Z. Shafiq, and F. Zaffar. Characterizing Seller-Driven Black-Hat Marketplaces. *arXiv preprint arXiv:1505.01637*, 2015.
- [24] M. Feldman, K. Lai, I. Stoica, and J. Chuang. Robust incentive techniques for peer-to-peer networks. In *Proceedings of the 5th ACM conference on Electronic commerce*, pages 102–111. ACM, 2004.
- [25] M. Feldman, C. Papadimitriou, J. Chuang, and I. Stoica. Free-riding and whitewashing in peer-to-peer systems. In *Proceedings of the ACM SIGCOMM workshop on Practice and theory of incentives in networked systems*, pages 228–236. ACM, 2004.
- [26] M. J. Freedman and R. Morris. Tarzan: A peer-to-peer anonymizing network layer. In *Proceedings of the 9th ACM conference on Computer and communications security*, pages 193–206. ACM, 2002.
- [27] M. Girvan and M. E. Newman. Community structure in social and biological networks. *Proceedings of the national academy of sciences*, 99(12):7821–7826, 2002.
- [28] C. Grier, K. Thomas, V. Paxson, and M. Zhang. @ spam: the underground on 140 characters or less. In *Proceedings of the 17th ACM conference on Computer and communications security*, pages 27–37. ACM, 2010.
- [29] D. Guilbeault and S. Woolley. How twitter bots are shaping the election. <http://www.theatlantic.com/technology/archive/2016/11/election-bots/506072/>, 11 2016.
- [30] M. Hearn. Modern anti-spam and e2e crypto. <https://moderncrypto.org/mail-archive/messaging/2014/000780.html>, 2014. Accessed: 2016-11-23.
- [31] E. Heilman, A. Kendler, A. Zohar, and S. Goldberg. Eclipse attacks on Bitcoin’s peer-to-peer network. In *24th USENIX Security Symposium (USENIX Security 15)*, pages 129–144, 2015.
- [32] Help Net Security. Twitter accounts spreading malicious code. <https://www.helpnetsecurity.com/2010/12/03/twitter-accounts-spreading-malicious-code/>, 12 2010. Accessed: 2016-11-2.
- [33] K. Hoffman, D. Zage, and C. Nita-Rotaru. A survey of attack and defense techniques for reputation systems. *ACM Computing Surveys (CSUR)*, 42(1):1, 2009.

- [34] T. Holz, M. Steiner, F. Dahl, E. Biersack, and F. C. Freiling. Measurements and Mitigation of Peer-to-Peer-based Botnets: A Case Study on Storm Worm. *LEET*, 8(1):1–9, 2008.
- [35] R. Housley, W. Polk, W. Ford, and D. Solo. Internet x. 509 public key infrastructure certificate and certificate revocation list (crl) profile. Technical report, 2002.
- [36] J. Jiang, Z.-F. Shan, X. Wang, L. Zhang, and Y.-F. Dai. Understanding Sybil Groups in the Wild. *Journal of Computer Science and Technology*, 30(6):1344–1357, 2015.
- [37] B. Jyothi and J. Dharanipragada. Symon: Defending large structured p2p systems against sybil attack. In *2009 IEEE Ninth International Conference on Peer-to-Peer Computing*, pages 21–30. IEEE, 2009.
- [38] S. D. Kamvar, M. T. Schlosser, and H. Garcia-Molina. The eigentrust algorithm for reputation management in p2p networks. In *Proceedings of the 12th international conference on World Wide Web*, pages 640–651. ACM, 2003.
- [39] C. Karlof and D. Wagner. Secure routing in wireless sensor networks: Attacks and countermeasures. *Ad hoc networks*, 1(2):293–315, 2003.
- [40] D. Koll, J. Li, J. Stein, and X. Fu. On the state of OSN-based Sybil defenses. In *Networking Conference, 2014 IFIP*, pages 1–9. IEEE, 2014.
- [41] E. Koutrouli and A. Tsalgatiidou. Taxonomy of attacks and defense mechanisms in P2P reputation systems—Lessons for reputation system designers. *Computer Science Review*, 6(2):47–70, 2012.
- [42] A. Kurve and G. Kesidis. Sybil detection via distributed sparse cut monitoring. In *2011 IEEE International Conference on Communications (ICC)*, pages 1–6. IEEE, 2011.
- [43] L. Lamport, R. Shostak, and M. Pease. The byzantine generals problem. *ACM Transactions on Programming Languages and Systems (TOPLAS)*, 4(3):382–401, 1982.
- [44] C. Lesniewski-Laas. A sybil-proof one-hop dht. In *Proceedings of the 1st workshop on Social network systems*, pages 19–24. ACM, 2008.
- [45] C. Lesniewski-Lass and M. F. Kaashoek. Whanau: A sybil-proof distributed hash table. NSDI, 2010.
- [46] J. Liang, N. Naoumov, and K. W. Ross. The Index Poisoning Attack in P2P File Sharing Systems. In *INFOCOM*, pages 1–12. Citeseer, 2006.
- [47] C. Liu, P. Gao, M. Wright, and P. Mittal. Exploiting temporal dynamics in Sybil defenses. In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security*, pages 805–816. ACM, 2015.
- [48] Y. Liu, S. Ji, and P. Mittal. Smartwalk: Enhancing social network security via adaptive random walks. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security, CCS '16*, pages 492–503, New York, NY, USA, 2016. ACM.
- [49] S. Marti and H. Garcia-Molina. Taxonomy of trust: Categorizing P2P reputation systems. *Computer Networks*, 50(4):472–484, 2006.
- [50] M. Meulpolder, J. A. Pouwelse, D. H. Epema, and H. J. Sips. Bartercast: A practical approach to prevent lazy freeriding in p2p networks. In *Parallel & Distributed Processing, 2009. IPDPS 2009. IEEE International Symposium on*, pages 1–8. IEEE, 2009.
- [51] A. Mislove, A. Post, P. Druschel, and P. K. Gummadi. Ostra: Leveraging Trust to Thwart Unwanted Communication. In *NSDI*, volume 8, pages 15–30, 2008.
- [52] A. Mislove, B. Viswanath, K. P. Gummadi, and P. Druschel. You are who you know: inferring user profiles in online social networks. In *Proceedings of the third ACM international conference on Web search and data mining*, pages 251–260. ACM, 2010.
- [53] S. Misra, A. S. M. Tayeen, and W. Xu. SybilExposer: An effective scheme to detect Sybil communities in online social networks. In *2016 IEEE International Conference on Communications (ICC)*, pages 1–6, May 2016.
- [54] A. Mohaisen and J. Kim. The Sybil attacks and defenses: a survey. *arXiv preprint arXiv:1312.6349*, 2013.
- [55] E. Mustafaraj and P. T. Metaxas. From obscurity to prominence in minutes: Political speech and real-time search. 2010.
- [56] S. Nakamoto. Bitcoin: A peer-to-peer electronic cash system, 2008.
- [57] A. Nandi, T.-W. J. Ngan, A. Singh, P. Druschel, and D. S. Wallach. Scrivener: Providing incentives in cooperative content distribution systems. In *Proceedings of the ACM/IFIP/USENIX 2005 International Conference on Middleware*, pages 270–291. Springer-Verlag New York, Inc., 2005.
- [58] J. Newsome, E. Shi, D. Song, and A. Perrig. The sybil attack in sensor networks: analysis & defenses. In *Proceedings of the 3rd international symposium on Information processing in sensor networks*, pages 259–268. ACM, 2004.
- [59] L. Page, S. Brin, R. Motwani, and T. Winograd. The PageRank citation ranking: bringing order to the web. 1999.
- [60] T. T. Project. Top-10 countries by directly connecting users. <https://metrics.torproject.org/userstats-relay-table.html>, 2016. Accessed: 2016-11-20.
- [61] G. Rakesh, S. Rangaswamy, V. Hegde, and G. Shoba. A survey of techniques to defend against sybil attacks in social networks. *International Journal of Advanced Research in Computer and Communication Engineering*, 3(5), 2014.
- [62] B. Ramsdell and S. Turner. Secure/multipurpose internet mail extensions (s/mime) version 3.2 message specification. Technical report, 2010.
- [63] J. Ratkiewicz, M. Conover, M. Meiss, B. Gonçalves, S. Patil, A. Flammini, and F. Menczer. Truthy: mapping the spread of astroturf in microblog streams. In *Proceedings of the 20th international conference companion on World wide web*, pages 249–252. ACM, 2011.
- [64] P. Resnick et al. The social cost of cheap pseudonyms. *Journal of Economics & Management Strategy*, 10(2):173–199, 2001.
- [65] SadBotTrue. Chapter 32. the stealth botnet, 6 2016. Accessed: 2016-11-2.

- [66] J.-M. Seigneur, A. Gray, and C. D. Jensen. Trust transfer: Encouraging self-recommendations without sybil attack. In *International Conference on Trust Management*, pages 321–337. Springer, 2005.
- [67] C. Selvaraj and S. Anand. A survey on security issues of reputation management systems for peer-to-peer networks. *Computer Science Review*, 6(4):145–160, 2012.
- [68] S. Seuken and D. C. Parkes. On the Sybil-proofness of accounting mechanisms. 2011.
- [69] S. Seuken and D. C. Parkes. Sybil-proof accounting mechanisms with transitive trust. In *Proceedings of the 2014 international conference on Autonomous agents and multi-agent systems*, pages 205–212. International Foundation for Autonomous Agents and Multiagent Systems, 2014.
- [70] L. Shi, S. Yu, W. Lou, and Y. T. Hou. Sybilshield: An agent-aided social network-based sybil defense among multiple communities. In *INFOCOM, 2013 Proceedings IEEE*, pages 1034–1042. IEEE, 2013.
- [71] K. C. Sia. DDoS vulnerability analysis of BitTorrent protocol. *UCLA: Technical Report*, 2006.
- [72] A. Singh et al. Eclipse attacks on overlay networks: Threats and defenses. In *In IEEE INFOCOM*. Citeseer, 2006.
- [73] Socialpuncher. How many primitive bots follow top-100? <http://socialpuncher.com/top-100/how-many-primitive-bots-follow-top-100/>, 9 2016. Accessed: 2016-11-2.
- [74] M. Srivatsa, L. Xiong, and L. Liu. TrustGuard: countering vulnerabilities in reputation management for decentralized overlay networks. In *Proceedings of the 14th international conference on World Wide Web*, pages 422–431. ACM, 2005.
- [75] T. Stein, E. Chen, and K. Mangla. Facebook immune system. In *Proceedings of the 4th Workshop on Social Network Systems*, page 8. ACM, 2011.
- [76] M. Steiner, T. En-Najjary, and E. W. Biersack. Exploiting KAD: possible uses and misuses. *ACM SIGCOMM Computer Communication Review*, 37(5):65–70, 2007.
- [77] G. Stringhini, C. Kruegel, and G. Vigna. Detecting spammers on social networks. In *Proceedings of the 26th Annual Computer Security Applications Conference*, pages 1–9. ACM, 2010.
- [78] D. Tamir. Twitter malware: Spreading more than just ideas. <https://securityintelligence.com/twitter-malware-spreading-more-than-just-ideas/>, 4 2013. Accessed: 2016-11-2.
- [79] D. N. Tran, B. Min, J. Li, and L. Subramanian. Sybil-Resilient Online Content Voting. In *NSDI*, volume 9, pages 15–28, 2009.
- [80] N. Tran, J. Li, L. Subramanian, and S. S. Chow. Optimal sybil-resilient node admission control. In *INFOCOM, 2011 Proceedings IEEE*, pages 3218–3226. IEEE, 2011.
- [81] B. Viswanath, A. Post, K. P. Gummadi, and A. Mislove. An analysis of social network-based sybil defenses. *ACM SIGCOMM Computer Communication Review*, 40(4):363–374, 2010.
- [82] L. Von Ahn, M. Blum, N. J. Hopper, and J. Langford. Captcha: Using hard ai problems for security. In *International Conference on the Theory and Applications of Cryptographic Techniques*, pages 294–311. Springer, 2003.
- [83] G. Wang, C. Wilson, X. Zhao, Y. Zhu, M. Mohanlal, H. Zheng, and B. Y. Zhao. Serf and turf: crowdturfing for fun and profit. In *Proceedings of the 21st international conference on World Wide Web*, pages 679–688. ACM, 2012.
- [84] L. Wang and J. Kangasharju. Real-world sybil attacks in BitTorrent mainline DHT. In *Global Communications Conference (GLOBECOM), 2012 IEEE*, pages 826–832. IEEE, 2012.
- [85] W. Wei, F. Xu, C. C. Tan, and Q. Li. Sybildefender: Defend against sybil attacks in large social networks. In *INFOCOM, 2012 Proceedings IEEE*, pages 1951–1959. IEEE, 2012.
- [86] J. Xue, Z. Yang, X. Yang, X. Wang, L. Chen, and Y. Dai. Votetrust: Leveraging friend invitation graph to defend against social network sybils. In *INFOCOM, 2013 Proceedings IEEE*, pages 2400–2408. IEEE, 2013.
- [87] C. Yang, R. Harkreader, J. Zhang, S. Shin, and G. Gu. Analyzing spammers’ social networks for fun and profit: a case study of cyber criminal ecosystem on twitter. In *Proceedings of the 21st international conference on World Wide Web*, pages 71–80. ACM, 2012.
- [88] M. Yang, Z. Zhang, X. Li, and Y. Dai. An empirical study of free-riding behavior in the maze p2p file-sharing system. In *International Workshop on Peer-to-Peer Systems*, pages 182–192. Springer, 2005.
- [89] H. Yu, P. B. Gibbons, M. Kaminsky, and F. Xiao. Sybillimit: A near-optimal social network defense against sybil attacks. In *2008 IEEE Symposium on Security and Privacy (sp 2008)*, pages 3–17. IEEE, 2008.
- [90] H. Yu, M. Kaminsky, P. B. Gibbons, and A. Flaxman. Sybilguard: defending against sybil attacks via social networks. In *ACM SIGCOMM Computer Communication Review*, volume 36, pages 267–278. ACM, 2006.
- [91] H. Yu, C. Shi, M. Kaminsky, P. B. Gibbons, and F. Xiao. Dsybil: Optimal sybil-resistance for recommendation systems. In *2009 30th IEEE Symposium on Security and Privacy*, pages 283–298. IEEE, 2009.